

PREDICTION MODEL 1

February 13, 2025

```
[1]: # Import necessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression, Lasso
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.metrics import (
    classification_report, confusion_matrix, roc_auc_score, accuracy_score,
    precision_score, recall_score, f1_score, roc_curve
)
from sklearn.decomposition import PCA
from sklearn.ensemble import VotingClassifier
from imblearn.over_sampling import SMOTE
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense

# 1. Load the dataset
print("Loading dataset...")
df = pd.read_csv('diabetes.csv')
print("\nDataset successfully loaded!")

# 2. Data Exploration
print("\nDataset Overview:")
print(df.head())
print("\nMissing Values:")
print(df.isnull().sum())
print("\nStatistical Summary:")
print(df.describe())

# Correlation heatmap
plt.figure(figsize=(10, 6))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm', fmt='.2f')
plt.title('Correlation Heatmap')
```

```

plt.show()

# Histograms for feature distributions
df.hist(figsize=(12, 10), bins=20, edgecolor='black')
plt.suptitle('Feature Distributions')
plt.show()

# Boxplots for key features
plt.figure(figsize=(10, 6))
sns.boxplot(data=df[['Glucose', 'BMI', 'Insulin']])
plt.title('Boxplot of Key Features')
plt.show()

# 3. Data Preprocessing
# Handling missing values
print("\nHandling missing values...")
df.fillna(df.mean(), inplace=True)
print("Missing values handled.\n")

# Splitting features and target
X = df.drop('Outcome', axis=1)
y = df['Outcome']

# Normalizing data
print("Normalizing data...")
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Addressing class imbalance using SMOTE
print("Addressing class imbalance using SMOTE...")
smote = SMOTE(random_state=42)
X_resampled, y_resampled = smote.fit_resample(X_scaled, y)
print("Class imbalance addressed.\n")

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(
    X_resampled, y_resampled, test_size=0.3, random_state=42
)

# 4. Dimensionality Reduction with PCA
print("Applying PCA...")
pca = PCA(n_components=2)
X_train_pca = pca.fit_transform(X_train)
X_test_pca = pca.transform(X_test)

# Visualizing PCA
plt.figure(figsize=(8, 6))

```

```

plt.scatter(X_train_pca[:, 0], X_train_pca[:, 1], c=y_train, cmap='coolwarm',
            edgecolors='k', alpha=0.7)
plt.title('PCA Visualization of Training Data')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.colorbar(label='Outcome')
plt.show()
print(f'Explained Variance Ratio by PCA Components: {pca.
      explained_variance_ratio_}\n')

# 5. Logistic Regression Model
print("Training Logistic Regression model...")
log_reg = LogisticRegression()
log_reg.fit(X_train, y_train)
y_pred_log_reg = log_reg.predict(X_test)

print("\nLogistic Regression Metrics:")
print(f"Accuracy: {accuracy_score(y_test, y_pred_log_reg)}")
print(f"Precision: {precision_score(y_test, y_pred_log_reg)}")
print(f"Recall: {recall_score(y_test, y_pred_log_reg)}")
print(f"F1 Score: {f1_score(y_test, y_pred_log_reg)}")
print(f"ROC-AUC: {roc_auc_score(y_test, log_reg.predict_proba(X_test)[: , 1])}")
print(confusion_matrix(y_test, y_pred_log_reg))

# 6. Decision Tree Model
print("\nTraining Decision Tree model...")
dt = DecisionTreeClassifier(max_depth=5, min_samples_split=5)
dt.fit(X_train, y_train)
y_pred_dt = dt.predict(X_test)

print("\nDecision Tree Metrics:")
print(f"Accuracy: {accuracy_score(y_test, y_pred_dt)}")
print(f"Precision: {precision_score(y_test, y_pred_dt)}")
print(f"Recall: {recall_score(y_test, y_pred_dt)}")
print(f"F1 Score: {f1_score(y_test, y_pred_dt)}")
print(f"ROC-AUC: {roc_auc_score(y_test, dt.predict_proba(X_test)[: , 1])}")
print(confusion_matrix(y_test, y_pred_dt))

# Visualizing the Decision Tree
plt.figure(figsize=(12, 8))
plot_tree(dt, filled=True, feature_names=df.columns[:-1],
          class_names=['Non-diabetic', 'Diabetic'], rounded=True)
plt.title('Decision Tree Visualization')
plt.show()

# 7. Neural Network Model
print("\nTraining Neural Network model...")

```

```

model = Sequential([
    Dense(16, input_dim=X_train.shape[1], activation='relu'),
    Dense(8, activation='relu'),
    Dense(1, activation='sigmoid')
])
model.compile(optimizer='adam', loss='binary_crossentropy',
    ↪metrics=['accuracy'])
history = model.fit(X_train, y_train, epochs=50, batch_size=10,
    ↪validation_data=(X_test, y_test), verbose=0)

# Neural Network Evaluation
nn_accuracy = model.evaluate(X_test, y_test, verbose=0)[1]
print(f"\nNeural Network Accuracy: {nn_accuracy}\n")

# 8. Ensemble Model: Voting Classifier
print("\nTraining Voting Classifier...")
voting_clf = VotingClassifier(estimators=[('log_reg', log_reg), ('dt', dt)],
    ↪voting='soft')
voting_clf.fit(X_train, y_train)
y_pred_voting = voting_clf.predict(X_test)

print("\nVoting Classifier Metrics:")
print(f"Accuracy: {accuracy_score(y_test, y_pred_voting)}")
print(f"ROC-AUC: {roc_auc_score(y_test, voting_clf.predict_proba(X_test)[:,-1])}")
print(classification_report(y_test, y_pred_voting))

# 9. Feature Importance Using Lasso Regularization
print("\nApplying Lasso Regularization...")
lasso = Lasso(alpha=0.1)
lasso.fit(X_train, y_train)
lasso_coefficients = pd.Series(lasso.coef_, index=df.columns[:-1])

# Plot Lasso coefficients
plt.figure(figsize=(10, 6))
lasso_coefficients.plot(kind='bar', color='skyblue')
plt.title('Lasso Regularization Feature Importance')
plt.ylabel('Coefficient Value')
plt.show()

# 10. Insights and Conclusions
print("\nKey Insights:")
print("1. Logistic Regression and Decision Tree models highlight Glucose, BMI,
    ↪and Insulin as key predictors.")
print("2. PCA visualization shows good separability between diabetic and
    ↪non-diabetic classes.")

```

```
print("3. Ensemble models provide improved accuracy and robustness.")
print("4. Lasso Regularization identifies the most important features for_
↳diabetes prediction.")
print("\nAnalysis complete!")
```

Loading dataset...

Dataset successfully loaded!

Dataset Overview:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI \
0	6	148	72	35	0	33.6
1	1	85	66	29	0	26.6
2	8	183	64	0	0	23.3
3	1	89	66	23	94	28.1
4	0	137	40	35	168	43.1

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

Missing Values:

Pregnancies	0
Glucose	0
BloodPressure	0
SkinThickness	0
Insulin	0
BMI	0
DiabetesPedigreeFunction	0
Age	0
Outcome	0


dtype: int64

Statistical Summary:

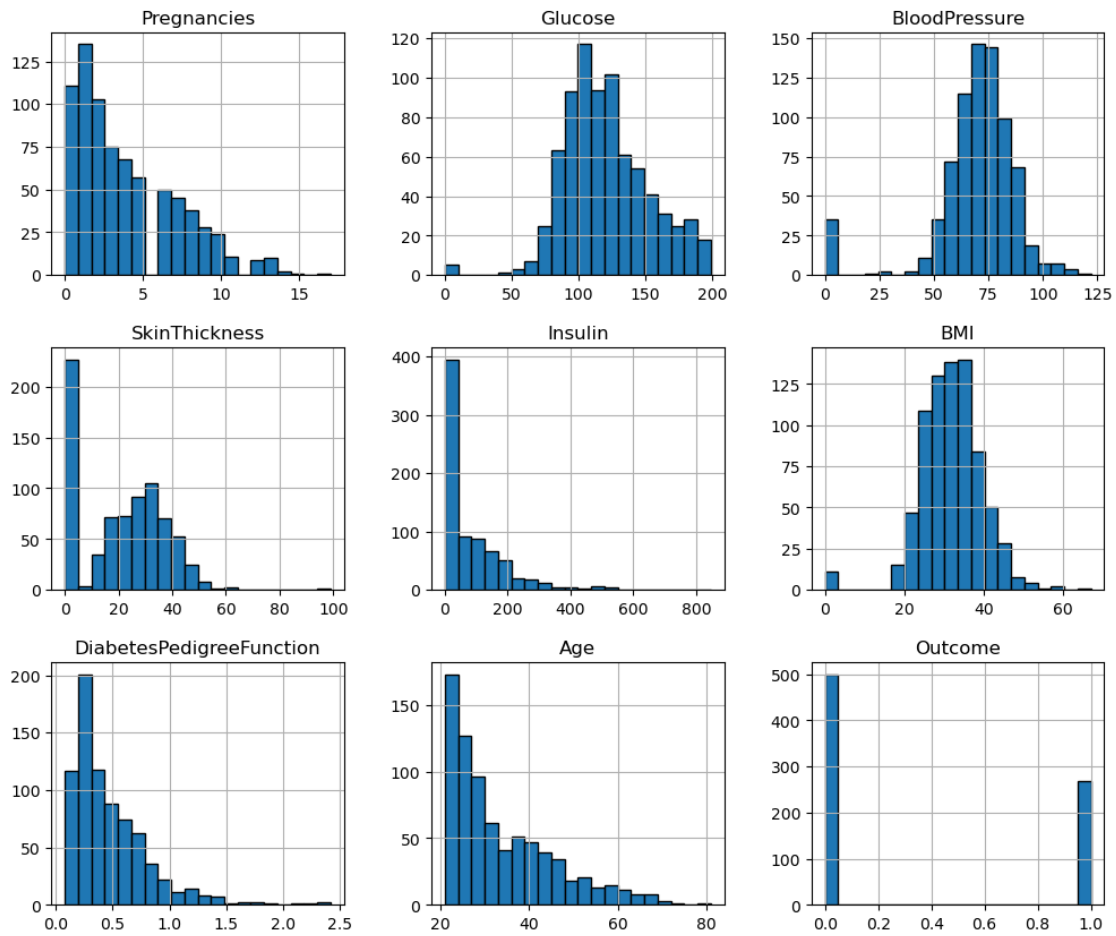
	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin \
count	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458	79.799479
std	3.369578	31.972618	19.355807	15.952218	115.244002
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	30.500000
75%	6.000000	140.250000	80.000000	32.000000	127.250000
max	17.000000	199.000000	122.000000	99.000000	846.000000


BMI	DiabetesPedigreeFunction	Age	Outcome
-----	--------------------------	-----	---------

count	768.000000	768.000000	768.000000	768.000000
mean	31.992578	0.471876	33.240885	0.348958
std	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.078000	21.000000	0.000000
25%	27.300000	0.243750	24.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000



Feature Distributions






Handling missing values...
Missing values handled.

Normalizing data...
Addressing class imbalance using SMOTE...
Class imbalance addressed.

Applying PCA...



Explained Variance Ratio by PCA Components: [0.26889982 0.21347626]

Training Logistic Regression model...

Logistic Regression Metrics:
Accuracy: 0.7533333333333333
Precision: 0.7655172413793103
Recall: 0.7350993377483444
F1 Score: 0.75
ROC-AUC: 0.8327036757189209
[[115 34]
[40 111]]

Training Decision Tree model...

Decision Tree Metrics:
Accuracy: 0.7366666666666667
Precision: 0.711764705882353
Recall: 0.8013245033112583

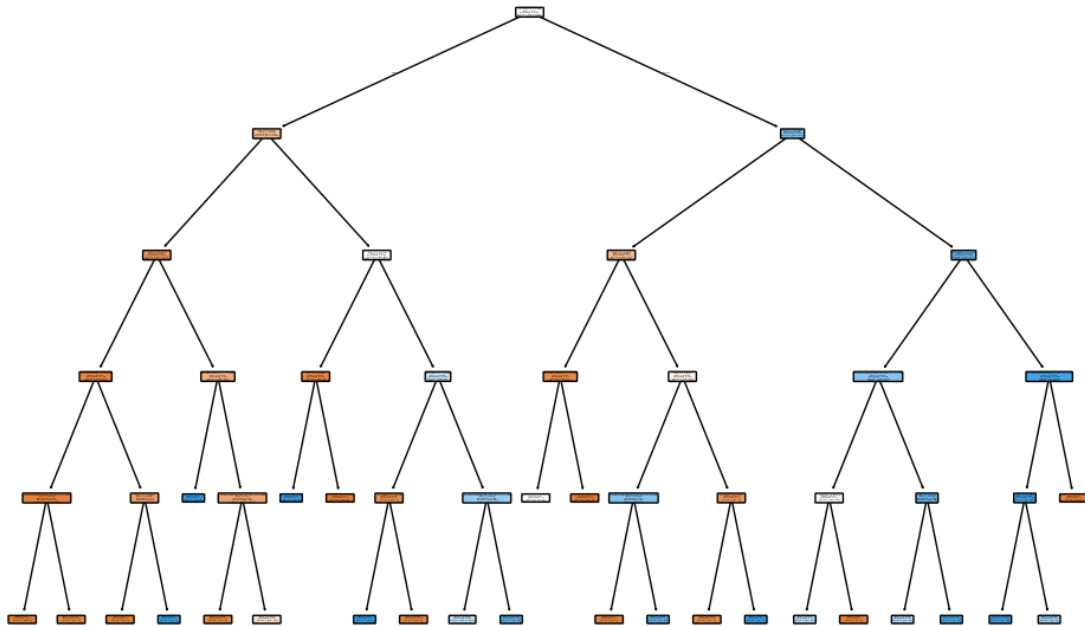
F1 Score: 0.7538940809968847

ROC-AUC: 0.804080181341393

[[100 49]

[30 121]]

Decision Tree Visualization



Training Neural Network model...

```
/opt/anaconda3/lib/python3.12/site-packages/keras/src/layers/core/dense.py:87:  
UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When  
using Sequential models, prefer using an `Input(shape)` object as the first  
layer in the model instead.
```

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

Neural Network Accuracy: 0.7633333206176758

Training Voting Classifier...

Voting Classifier Metrics:

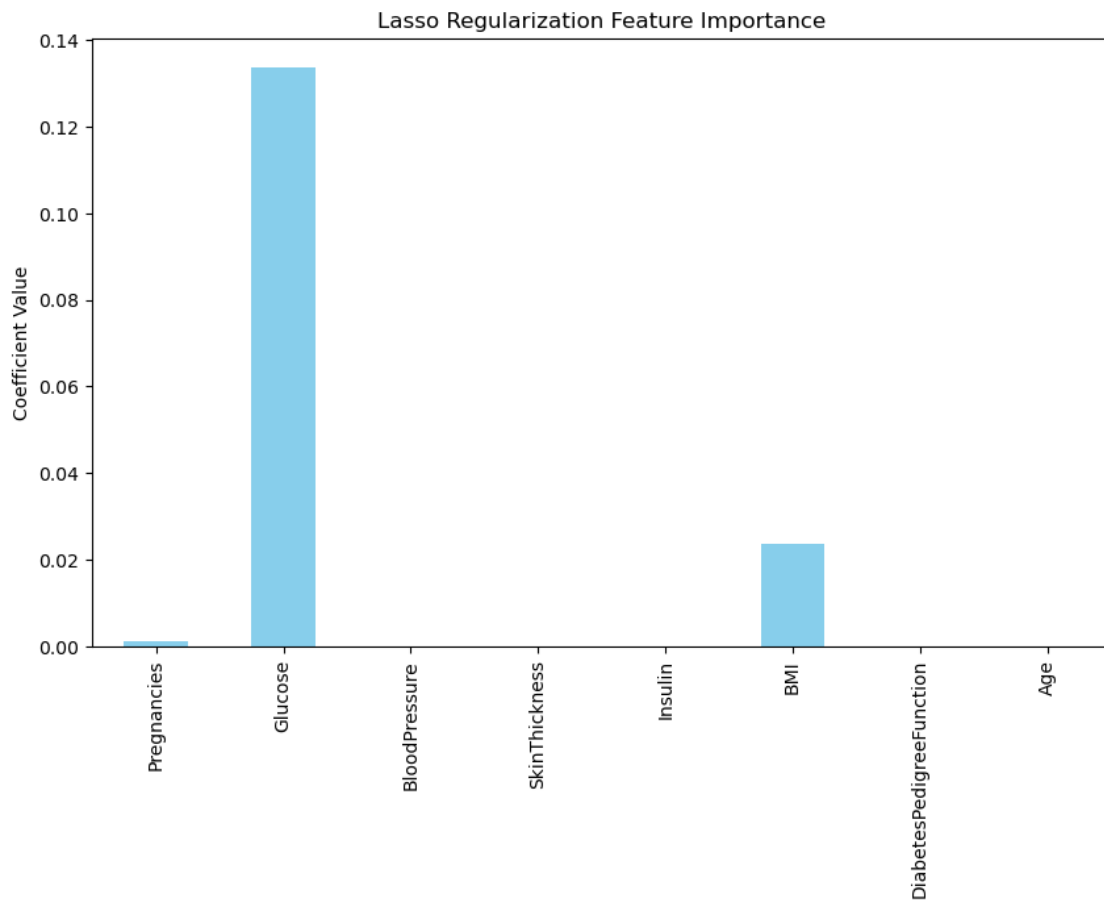
Accuracy: 0.7466666666666667

ROC-AUC: 0.8447931019156407

precision	recall	f1-score	support
-----------	--------	----------	---------

	0	0.75	0.73	0.74	149
	1	0.74	0.76	0.75	151
accuracy				0.75	300
macro avg		0.75	0.75	0.75	300
weighted avg		0.75	0.75	0.75	300

Applying Lasso Regularization...



Key Insights:

1. Logistic Regression and Decision Tree models highlight Glucose, BMI, and Insulin as key predictors.
2. PCA visualization shows good separability between diabetic and non-diabetic classes.
3. Ensemble models provide improved accuracy and robustness.
4. Lasso Regularization identifies the most important features for diabetes

prediction.

Analysis complete!